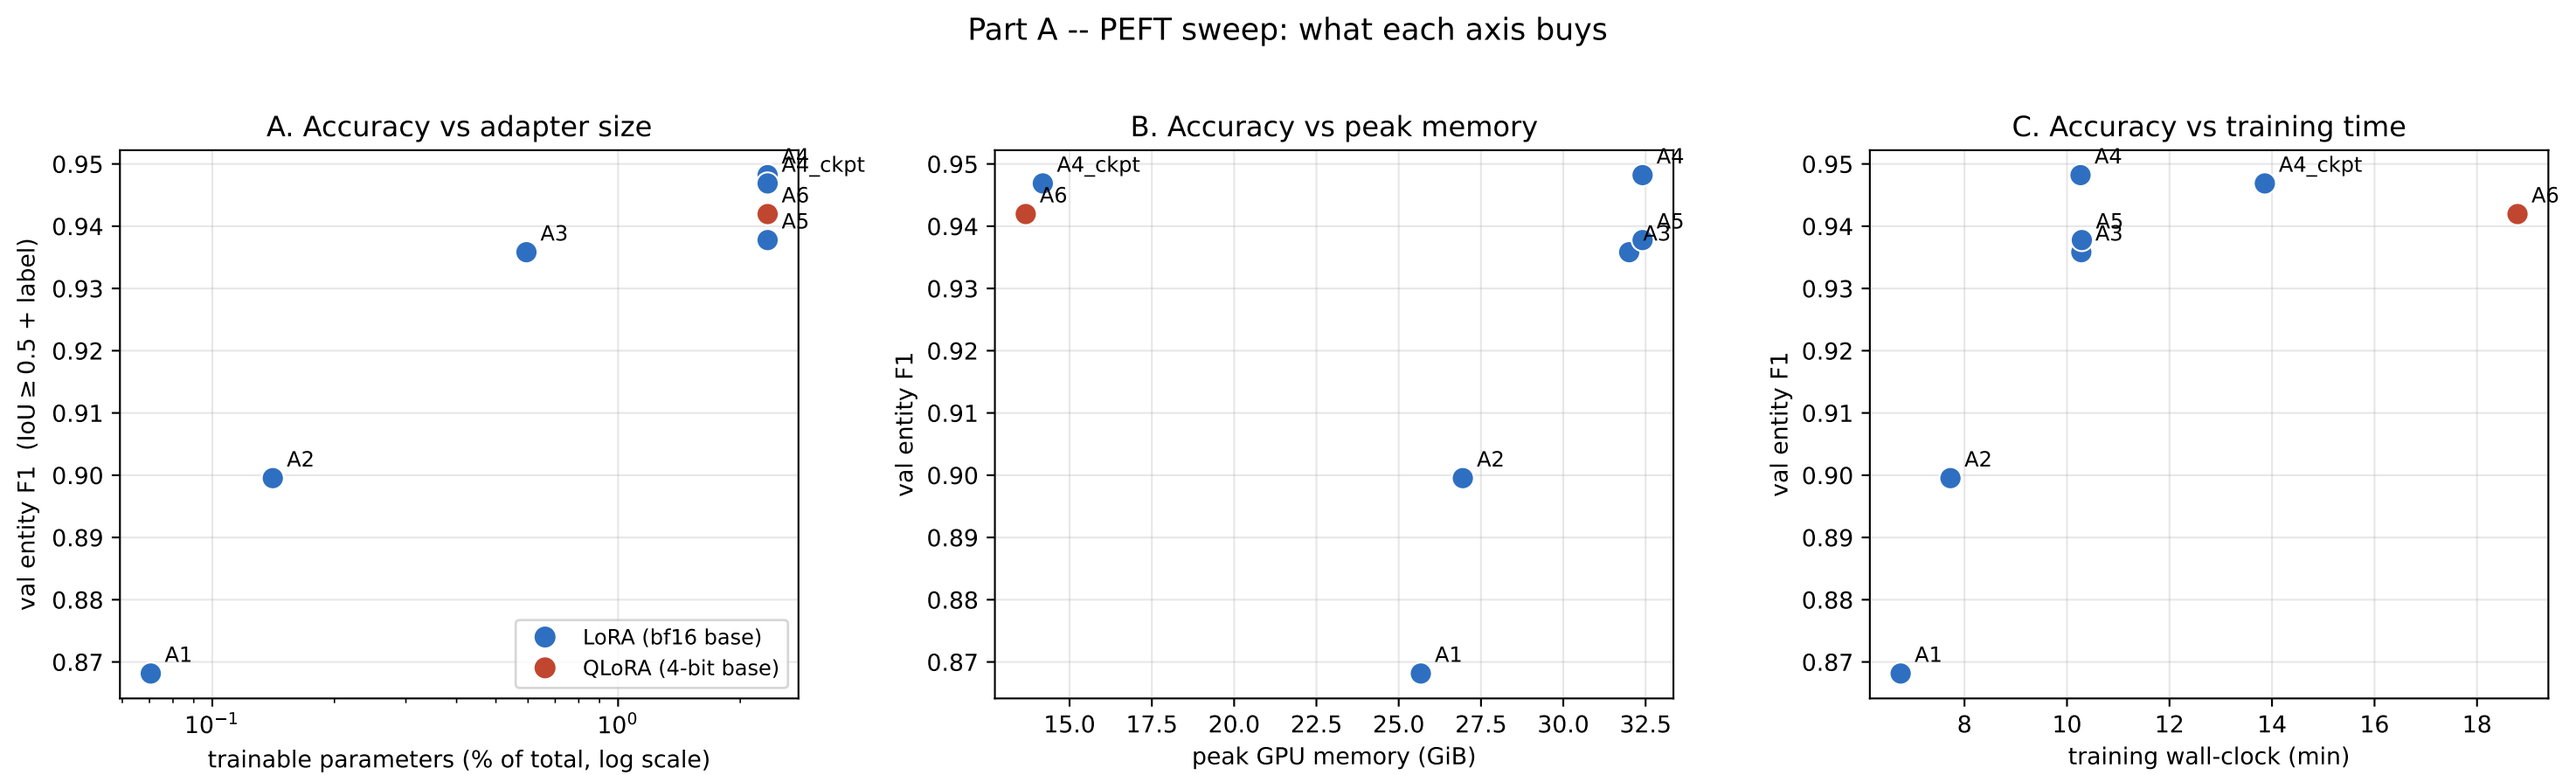
Part A -- PEFT sweep: what each axis buys
A. Accuracy vs adapter size
val entity F1 (IoU ≥ 0.5 + label)
trainable parameters (% of total, log scale)
A1
A2
A3
A4
A4_ckpt
A5
A6
LoRA (bf16 base)
QLoRA (4-bit base)
B. Accuracy vs peak memory
val entity F1
peak GPU memory (GiB)
C. Accuracy vs training time
val entity F1
training wall-clock (min)